

Day 2, Session 1: Introduction to Pandas

Three-Day Data Analysis with Python Course

Yesterday's Recap

We covered Python fundamentals:

- Variables, data types, operators
- Control flow (if/else, loops)
- Functions
- Data structures (lists, dicts, tuples, sets)
- NumPy arrays
- File operations

Today: We move from Python basics to real data analysis!

Overview of Session 1

1. From SQL to Pandas
2. Pandas Fundamentals
3. Data Loading and Inspection
4. Data Selection and Indexing

Part 1: From SQL to Pandas

Why Pandas?

pandas is the industry-standard library for data analysis in Python.

- Built on top of NumPy
- Powerful data manipulation capabilities
- Intuitive syntax similar to SQL
- Handles messy real-world data
- Integrates seamlessly with visualization libraries

If you know SQL, you'll feel right at home!

SQL Tables vs DataFrames

Concept	SQL	Pandas
Data structure	Table	DataFrame
Row	Record/Row	Row
Column	Column/Field	Column/Series
Filter rows	<code>WHERE</code>	Boolean indexing
Select columns	<code>SELECT</code>	Column selection
Aggregate	<code>GROUP BY</code>	<code>.groupby()</code>
Join tables	<code>JOIN</code>	<code>.merge()</code>
Sort	<code>ORDER BY</code>	<code>.sort_values()</code>

Row-Based vs Columnar Storage

SQL (Row-based):

- Optimized for transactional operations
- Reads entire rows at a time

Pandas (Columnar):

- Optimized for analytical operations
- Reads columns at a time
- Much faster for aggregations and calculations

Why DataFrames Are Powerful

Key advantages:

1. **Labeled axes** - Rows and columns have names
2. **Mixed types** - Different columns can have different types
3. **Size mutable** - Can add/remove rows and columns
4. **Missing data** - Built-in handling of missing values
5. **Group operations** - Easy split-apply-combine
6. **Alignment** - Automatic data alignment by label

Part 2: Pandas Fundamentals

Importing Pandas

```
# Import pandas (standard alias is 'pd')
import pandas as pd

# Also import numpy (often used together)
import numpy as np

print(f"pandas version: {pd.__version__}")
```

Convention: Always use `pd` as the alias for pandas.

Series: One-Dimensional Labeled Arrays

A **Series** is a one-dimensional array with labels (index).

- Similar to a column in a spreadsheet
- Built on top of NumPy arrays
- Has an index and values

Syntax:

```
series = pd.Series(data, index=index)
```

Series Examples

```
# Create a Series from a list
sales = pd.Series([1200, 1450, 1100, 1800, 2100])
print(sales)

# Create a Series with custom index
daily_sales = pd.Series(
    [1200, 1450, 1100, 1800, 2100],
    index=['Mon', 'Tue', 'Wed', 'Thu', 'Fri']
)
print(daily_sales)

# Access by index
print(f"Monday sales: ${daily_sales['Mon']}")
```

Series (Continued)

```
# Series from dictionary
prices = pd.Series({
    'Laptop': 999.99,
    'Mouse': 29.99,
    'Keyboard': 79.99,
    'Monitor': 299.99
})
print(prices)

# Series have attributes
print(f"Index: {prices.index}")
print(f"Values: {prices.values}")
print(f>Data type: {prices.dtype}")
```

DataFrames: Two-Dimensional Labeled Data

A **DataFrame** is a two-dimensional labeled data structure.

- Think of it as a table or spreadsheet
- Collection of Series (each column is a Series)
- Has both row index and column names

This is the primary data structure you'll work with!

DataFrame Anatomy

	Column1	Column2	Column3
Index1	val11	val12	val13
Index2	val21	val22	val23
Index3	val31	val32	val33

- **Columns:** Named Series
- **Index:** Row labels (like primary key)
- **Values:** The actual data

Creating DataFrames

Multiple ways to create DataFrames:

1. From dictionaries
2. From lists of dictionaries
3. From NumPy arrays
4. **From files** (most common!)

Let's see examples...

DataFrame from Dictionary

```
# Dictionary of lists (each key becomes a column)
data = {
    'product': ['Laptop', 'Mouse', 'Keyboard', 'Monitor'],
    'price': [999.99, 29.99, 79.99, 299.99],
    'stock': [15, 120, 45, 30],
    'category': ['Computer', 'Accessory', 'Accessory', 'Computer']
}

df = pd.DataFrame(data)
print(df)
```

DataFrame from List of Dictionaries

```
# List of dictionaries (each dict becomes a row)
customers = [
    {'id': 101, 'name': 'Alice', 'city': 'NYC', 'purchases': 15},
    {'id': 102, 'name': 'Bob', 'city': 'LA', 'purchases': 8},
    {'id': 103, 'name': 'Carol', 'city': 'Chicago', 'purchases': 22}
]

customers_df = pd.DataFrame(customers)
print(customers_df)
```

Part 3: Data Loading and Inspection

Reading CSV Files

The most common way to create DataFrames:

```
# Read CSV file
df = pd.read_csv('filename.csv')

# Common parameters
df = pd.read_csv(
    'filename.csv',
    sep=',',          # Delimiter (default is comma, sometimes "\t")
    header=0,         # Row number for column names
    index_col=None,    # Column to use as index
    na_values=['NA']   # Additional strings to recognize as NaN
)
```

Reading Other File Formats

Pandas supports many file formats:

```
# JSON files
df = pd.read_json('data.json')

# Excel files
df = pd.read_excel('data.xlsx', sheet_name='Sheet1')

# SQL databases
df = pd.read_sql('SELECT * FROM table', connection)

# Many more: HDF5, Parquet...
```

Example: Reading Customer Data

```
# Let's load some customer data
customers = pd.read_csv('customers.csv')

# What do we have?
print(type(customers)) # DataFrame
print(f"Shape: {customers.shape}") # (rows, columns)
```

Basic DataFrame Inspection

Method	Description
<code>.head(n)</code>	First n rows (default 5)
<code>.tail(n)</code>	Last n rows (default 5)
<code>.info()</code>	Column types and missing values
<code>.describe()</code>	Statistical summary
<code>.shape</code>	Dimensions (rows, columns)
<code>.columns</code>	Column names
<code>.dtypes</code>	Data types of each column

.head() and **.tail()**

```
# View first 5 rows
print(customers.head())

# View first 3 rows
print(customers.head(3))

# View last 5 rows
print(customers.tail())
```

Always start with `.head()` to see what you're working with!

`.info()` - Understanding Your Data

```
customers.info()
```

Output:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100 entries, 0 to 99
Data columns (total 5 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   id          100 non-null   int64
 1   name        100 non-null   object
 2   email       98 non-null    object
 3   purchases  100 non-null   int64
 4   total       100 non-null   float64
dtypes: float64(1), int64(2), object(2)
memory usage: 4.0+ KB
```

`.describe()` - Statistical Summary

```
customers.describe()
```

Output:

	id	purchases	total
count	100.00000	100.000000	100.000000
mean	150.50000	12.450000	1845.234000
std	29.01149	7.823421	892.451234
min	101.00000	1.000000	150.500000
25%	125.75000	6.000000	1021.750000
50%	150.50000	11.500000	1789.250000
75%	175.25000	18.000000	2543.000000
max	200.00000	35.000000	4950.000000

Only includes numerical columns by default!

Understanding Data Types

Common pandas dtypes:

Dtype	Python Type	Usage
<code>int64</code>	<code>int</code>	Integer numbers
<code>float64</code>	<code>float</code>	Decimal numbers
<code>object</code>	<code>str</code>	Text/strings
<code>bool</code>	<code>bool</code>	True/False
<code>datetime64</code>	<code>datetime</code>	Dates and times
<code>category</code>	-	Categorical data

Shape and Size

```
# Get dimensions
rows, cols = customers.shape
print(f"Rows: {rows}, Columns: {cols}")

# Number of elements
print(f"Total elements: {customers.size}")

# Number of rows
print(f"Row count: {len(customers)}")
```

Part 4: Data Selection and Indexing

Selecting Columns

Several ways to select columns:

```
# Single column (returns Series)
names = customers['name']
print(type(names)) # Series

# Single column alternative (if no spaces in name)
names = customers.name

# Multiple columns (returns DataFrame)
subset = customers[['name', 'email']]
print(type(subset)) # DataFrame
```

Note: Double brackets `[[...]]` for multiple columns!

Selecting Columns (Example)

```
# Select single column
print("Customer names:")
print(customers['name'].head())

# Select multiple columns
print("\nCustomer contact info:")
print(customers[['name', 'email']].head())

# Calculate on a column
total_purchases = customers['purchases'].sum()
print(f"\nTotal purchases: {total_purchases}")
```

Selecting Rows by Position

Use `.iloc[]` for integer-location based indexing:

```
# First row
first_customer = customers.iloc[0]

# First 5 rows
first_five = customers.iloc[0:5]

# Last row
last_customer = customers.iloc[-1]

# Specific rows
some_rows = customers.iloc[[0, 5, 10]]
```


Selecting Rows by Label

Use `.loc[]` for label-based indexing:

```
# If index is set to a column (e.g., customer ID)
customers_indexed = customers.set_index('id')

# Select by index label
customer_125 = customers_indexed.loc[125]

# Select range of labels
customers_range = customers_indexed.loc[101:105]
```

`.loc` vs `.iloc`

Method	Type	Example
<code>.iloc[n]</code>	Integer position	<code>df.iloc[0]</code> (first row)
<code>.loc[label]</code>	Index label	<code>df.loc['2024-01-01']</code>

Key difference:

- `.iloc` - Think "i" for "integer"
- `.loc` - Think "l" for "label"

Boolean Indexing (Filtering)

The most powerful selection method - like SQL's `WHERE` clause!

```
# Filter rows where purchases > 10
active_customers = customers[customers['purchases'] > 10]

# Filter with multiple conditions (use & and |)
premium = customers[
    (customers['purchases'] > 15) &
    (customers['total'] > 2000)
]

# Filter by string matching
nyc_customers = customers[customers['city'] == 'NYC']
```

This is how you'll filter data most of the time!

Boolean Indexing (Example)

```
# High-value customers
high_value = customers[customers['total'] > 2000]
print(f"High-value customers: {len(high_value)}")

# Active customers in NYC
active_nyc = customers[
    (customers['purchases'] > 10) &
    (customers['city'] == 'NYC')
]
print(f"\nActive NYC customers:\n{active_nyc[['name', 'purchases']]}")
```

Combining Selection Methods

You can combine row and column selection:

```
# Select specific rows and columns with .loc
subset = customers.loc[0:5, ['name', 'total']]

# Boolean filter + column selection
high_value_names = customers.loc[
    customers['total'] > 2000,
    ['name', 'total']
]

# .iloc with row and column positions
subset = customers.iloc[0:10, 0:3]
```

Practical Example: Customer Segmentation

```
# Load customer data
customers = pd.read_csv('customers.csv')

# View structure
print(customers.info())
print(customers.head())

# Segment 1: VIP customers (>20 purchases, >$3000 total)
vip = customers[
    (customers['purchases'] > 20) &
    (customers['total'] > 3000)
]

# Segment 2: At-risk (few purchases, low total)
at_risk = customers[
    (customers['purchases'] < 5) &
    (customers['total'] < 500)
]
```

Practical Example (Continued)

```
# Segment 3: Active (10-20 purchases)
active = customers[
    (customers['purchases'] >= 10) &
    (customers['purchases'] <= 20)
]

# Analyze segments
print(f"VIP customers: {len(vip)} ({len(vip)/len(customers)*100:.1f}%)")
print(f"At-risk customers: {len(at_risk)} ({len(at_risk)/len(customers)*100:.1f}%)")
print(f"Active customers: {len(active)} ({len(active)/len(customers)*100:.1f}%)")

# Average spending by segment
print(f"\nAverage VIP spending: ${vip['total'].mean():.2f}")
print(f"Average at-risk spending: ${at_risk['total'].mean():.2f}")
```

Session 1 Summary

- ✓ **SQL to Pandas:** DataFrames are like SQL tables with superpowers
- ✓ **Fundamentals:** Series and DataFrames are the core structures
- ✓ **Loading Data:** `read_csv()`, `read_excel()`, and other readers
- ✓ **Inspection:** `.head()`, `.info()`, `.describe()`, `.shape`
- ✓ **Selection:** Column selection, `.loc`, `.iloc`, boolean indexing

Questions?

Break time! 

See you in Session 2!